
Re-Implementing CNN Price-Trend Signals: A Reproduction of Jiang, Kelly, and Xiu (2023)

Kaibiao Zhu

Final Project, Machine Learning for Finance
The Hong Kong University of Science and Technology (Guangzhou)
kzhu597@connect.hkust-gz.edu.cn

Abstract

This paper reproduces the image-based price-trend prediction framework of Jiang et al. [2023] (hereafter the *original paper*), which replaces predefined technical indicators with OHLC-volume charts and convolutional neural networks (CNNs). Using CRSP daily U.S. equity data, we reconstruct the image pipeline, train horizon-specific CNN ensembles, and evaluate the resulting signals through out-of-sample backtests. We compare the CNN signals with momentum, moving-average, and reversal benchmarks, and assess residual performance through factor-spanning regressions. The reproduction is strongest at short horizons, where image-based signals deliver the best economic performance. Relative to the magnitudes reported by the original paper, medium- and long-horizon strategies weaken substantially after costs, and factor stripping yields significantly negative alphas. Saliency maps indicate that the CNN mainly attends to recent price movement and volume contrast. We also provide replication materials so that the empirical results can be audited and extended.

1 Introduction

Technical analysis has long suggested that the visual “shape” of a price chart conveys information beyond what is captured by scalar momentum or moving-average filters. The original paper formalises this intuition by feeding fixed-resolution open–high–low–close (OHLC) images with a volume sub-panel into a convolutional neural network and showing that the resulting forecasts dominate canonical price-trend signals out of sample. The study joins a growing empirical asset-pricing literature that uses flexible machine-learning predictors to map raw historical features into expected returns [Gu et al., 2020, Kelly et al., 2019], and it ties directly to early efforts at formalising chart patterns via kernel smoothing [Lo et al., 2000].

This paper re-implements the core pipeline from end to end and evaluates it under the same in-sample and out-of-sample calendar design as the original paper: training and validation draw from 1993–2000 with a 70/30 random sample split, and headline out-of-sample results are reported for 2001–2019. Our contributions are:

1. An independent OHLC-plus-volume image renderer and horizon-matched CNN stacks built from the dimensions, drawing rules, and layer geometry documented in the original paper—including the simple moving-average overlay, volume sub-panel, and vertical orientation convention—implemented from the original paper’s written specification and reported tensor dimensions.
2. A CRSP-style daily panel for data construction, with headline out-of-sample CNN evaluation as above; proportional transaction costs; matched-horizon price momentum, simple moving-average cross, and short-term reversal baselines; and a five-factor-plus-momentum spanning regression [Fama and French, 1993, 2015, Carhart, 1997].

3. Input-gradient saliency analysis [Simonyan et al., 2014] of the $I = 60$ ensemble, which highlights the right tail of each image and the volume bars as the dominant feature regions.
4. Code, trained weights, and post-aggregation CSV tables are released on GitHub, allowing every numerical claim to be regenerated.

The remainder of the paper is organised as follows. Section 2 positions our reproduction relative to the chart-pattern, deep-learning, and interpretability literatures. Section 3 details data sources and image construction. Section 4 describes the CNN architecture and training recipe. Section 5 reports empirical results (decile prediction accuracy, baseline comparisons, cumulative returns, classification metrics, decile portfolio statistics, and Fama–French spanning regressions). Section 6 visualises model attention. Section 7 discusses where our results align with and diverge from the original paper.

2 Related work

Chart-pattern literature. Lo et al. [2000] provided one of the first systematic treatments of technical chart patterns by combining kernel regression with automated pattern recognition. Jegadeesh and Titman [1993] documented price momentum as a robust cross-sectional return predictor, providing the canonical baseline against which chart-based technical trend signals are often benchmarked empirically.

Deep learning for return prediction. Recent work demonstrates that deep architectures uncover return signals that are difficult to capture with closed-form characteristics. Gu et al. [2020] benchmark a broad set of supervised learners on U.S. equities; Kelly et al. [2019] use latent-factor models with characteristic loadings. Within technical analysis, Sezer and Ozbayoglu [2018] convert indicators into images for trading. The original paper unifies these threads by treating the OHLC chart itself as the model input.

Interpretability via input gradients. Simonyan et al. [2014] introduced gradient-based saliency to visualise which input pixels most influence a CNN’s class score; we apply the same technique to identify which segments of each OHLC image drive the model’s “up” probability. CNN architectures used here trace back to LeCun et al. [1998].

3 Data and image construction

Universe. We employ a CRSP-consistent daily security-level panel of U.S. equities spanning 1992-01-01, to 2024-12-31. Data are observed at the daily frequency for each security, uniquely identified by its CRSP permanent security number (PERMNO). The panel includes core trading and pricing fields: daily market capitalization, total and price returns, trading volume, and open-high-low-close (OHLC) prices. We impose a standard sequential screening procedure: all fields are required to be non-missing; we retain only securities with a closing price of at least \$1 and strictly positive trading volume. Log returns are computed, and OHLC prices are adjusted for corporate actions via cumulative log-return compounding, with the first valid observation for each security serving as the normalization anchor. No additional survivorship bias adjustments are applied beyond those inherent to the standard CRSP universe construction.

Labels and holding period. For each horizon $F \in \{5, 20, 60\}$ trading days we form $R_{t,F}^{\text{fut}} = \exp\left(\sum_{s=t+1}^{t+F} \log(1 + r_s)\right) - 1$ on the adjusted price path, and a binary classification label $y_{t,F} = \mathbf{1}\{R_{t,F}^{\text{fut}} > 0\}$. Rebalance timestamps are fixed at period-end calendar boundaries (Friday for $F = 5$, month-end for $F = 20$, quarter-end for $F = 60$), matching the canonical SPY trade calendar.

Image construction. Following the original paper, we encode each stock–date observation as a fixed-size grayscale *chart image* rather than a vector of technical indicators. The background is black (pixel value 0) and all price–volume objects are drawn in white (255). We consider three input horizons $I \in \{5, 20, 60\}$, each matched to the prediction horizon $F = I$, so the image summarizes the most recent I trading days of adjusted OHLC prices and daily volume.

OHLC panel. The upper portion of the image is reserved for open–high–low–close prices. Each trading day τ in the I -day window occupies three adjacent pixel columns: a one-pixel open marker in the left column, a vertical high–low bar in the centre, and a one-pixel close marker in the right column. Prices are linearly mapped into the vertical axis using the window minimum and maximum, where the range is defined to include both the raw OHLC values and the moving-average line described below. After drawing, the OHLC sub-panel is flipped top-to-bottom so that upward price movements read visually as “up” on the chart, in line with conventional price plots.

Moving-average line. We overlay a simple moving average (SMA) of the close price, using an I -day window—the same length as the chart horizon (e.g., a 20-day chart uses a 20-day SMA). For each day, the SMA value is placed on the centre column of that day; consecutive days are connected by a one-pixel poly-line. To ensure that the first in-window SMA already reflects I days of history, the SMA is computed from $2I - 1$ days of past closes ending at the chart date.

Volume sub-panel. The lower portion of the image displays volume bars. As in the original paper, the volume chart is drawn along the bottom of the image and occupies one-fifth of the total height; the OHLC panel fills the upper four-fifths, with a one-pixel gap between the two regions. Within the volume band, the largest daily volume in the window is scaled to the top of the band and all other days are drawn proportionally. Missing-volume days are left blank.

Image dimensions. Because each day uses three pixel columns, the image width is $W = 3I$. The total height H stacks the OHLC block, the separator, and the volume band. Table 1 reports the (W, H) pairs used in our implementation, identical to those in the original paper.

Table 1: Image Dimensions

Horizon	W	H
$I = 5$	15	32
$I = 20$	60	64
$I = 60$	180	96

Figure 1 provides a concrete illustration of the construction rules above. We plot the same security on the same calendar anchor for $I \in \{5, 20, 60\}$, so the only deliberate difference across panels is the look-back length. The left panel ($I = 5$) is the coarsest view: fifteen columns of price action and a compact volume strip. The centre and right panels progressively widen the canvas to sixty and one hundred eighty columns, revealing more fine-grained OHLC structure, a longer SMA trace, and richer intra-window variation in volume. Because every image fed to the CNN has fixed (W, H) for a given horizon, longer windows do not add pixels per trading day; instead, the model sees the same economic history at three standardised resolutions, analogous to viewing a price chart at weekly, monthly, and quarterly zoom levels.




Figure 1: OHLC-plus-volume images for a single security at horizons $I \in \{5, 20, 60\}$.

Train/validation/test split. We follow the original paper in treating the early portion of the panel as in-sample. Training is performed on 1993–2000 with a 70/30 random sample-level split for training/validation. Out-of-sample backtests start in January 2001. Headline tables and figures report results through December 2019, matching the evaluation window in the original paper.

4 CNN architecture and training

Architectures. The three horizon-specific backbones mirror the published configuration. All blocks use 5×3 kernels, a LeakyReLU(0.01) activation, batch normalisation, and $(2, 1)$ MaxPool with `ceil_mode=True` (the $I = 5$ head uses the standard floor mode). $I = 20$ and $I = 60$ insert dilated and strided first convolutions so the receptive field can span longer windows without an explosion in parameter count. Figure 2 summarises the three stacks; layer-level details are in Table 5 (Appendix B).




Figure 2: Horizon-specific CNN backbones (cf. Figure 2 of the original paper). Each rectangle is one block: 5×3 convolution (with the channel width shown), LeakyReLU, batch normalisation, and 2×1 max pooling. Stride and dilation for $I = 20$ and $I = 60$ are in Table 5. FC widths 15,360 / 46,080 / 184,320 match our implementation.

Optimisation. We minimise binary cross-entropy on the per-horizon label $y_{t,F}$ using stochastic gradient descent (Adam variant disabled), with learning rate 10^{-5} , batch size 128, and a maximum of 50 epochs with early-stopping on validation loss. Convolutional and linear weights are initialised with Xavier-uniform and biases at zero. We train an ensemble of five seeds $\{42, 123, 456, 789, 101112\}$ and aggregate by averaging the softmax-probability outputs at inference time. The resulting model sizes are approximately 0.6 MB ($I = 5$), 2.8 MB ($I = 20$), and 11.8 MB ($I = 60$) per seed, consistent with the parameter counts implied by Figure 2.

5 Empirical results

We score every test-set observation with the five-seed ensemble and form ten cross-sectional deciles by predicted “up”-probability at each rebalance date. The high-minus-low (H–L) portfolio is equal-weighted within deciles, rebalanced at horizon-end, and adjusted for proportional transaction costs at 10 bp per side (20 bp round-trip) applied to both legs.

5.1 Decile prediction accuracy (Figure 6 reproduction)

At $I = 5$ the CNN decile annualised returns rise nearly monotonically from 6.62% in D1 to 20.67% in D10 (Figure 3a), while the five-day momentum and moving-average curves are essentially flat across the cross-section. At longer horizons the picture changes: the $I = 20$ CNN spread between D10 and D1 is mildly positive (11.10% vs. 9.78%; Figure 3b), and at $I = 60$ the CNN ordering is essentially flat (11.95% in D10 against 13.00% in D1; Figure 3c). The traditional baselines, by contrast, display a U-shaped pattern at $I = 20$ and $I = 60$: both the bottom and top deciles deliver above-average returns, reflecting the well-known concentration of momentum and mean-reversion premia in the extreme deciles.




Figure 3: Annualised decile returns by horizon for the CNN, matched-horizon momentum, and moving-average cross signals. Reproduction of Figure 6 of the original paper; all series are computed from the same out-of-sample prediction panels and rebalance calendar.




Figure 4: Simplified Figure 7 reproduction: CNN deciles against momentum, moving-average cross, and short-term reversal baselines. At $I = 5$ the reversal baseline’s D10 reaches an annualised return of 37.78%, against which the CNN’s monotone ranking is more useful as a signal than as a stand-alone strategy.

5.2 Comparison with traditional baselines (Figure 7)

Figure 4 shows that at $I = 5$ the CNN ranking strictly dominates both momentum and moving-average cross across the cross-section, but is itself dominated in absolute level by the short-term reversal baseline whose D10 annualised return reaches 37.78% (driven by the microstructure-style bid–ask bounce on which it relies). At $I = 20$ the momentum and moving-average curves display the same U-shaped pattern noted above, and the CNN’s monotone ranking is qualitatively closer to the monthly-frequency pattern the original paper emphasises when comparing CNN deciles to traditional signals. At $I = 60$ the CNN ranking is nearly horizontal both here and in Figure 6(c) of the original paper, consistent with the well-documented attenuation of price-trend signals at quarterly horizons.

5.3 Cumulative returns

Figure 5 plots the cumulative log return of the H–L spread. The $I = 5$ net curve climbs steadily because the high-frequency ranking is sharp enough to absorb the transaction-cost drag. At $I = 20$ and $I = 60$, the gross curves are mildly positive but the net curves are substantially attenuated: with turnover at 84.08% and 86.68% per period (Table 2), the 20 bp round-trip cost erases most of the gross spread.

5.4 Decile portfolio statistics

Table 2 confirms the pattern visible in Figures 3–5: only the $I = 5$ horizon delivers a robust net H–L spread (annualised SR 1.29). At $I = 20$ the H–L SR is 0.16, and at $I = 60$ the spread is mildly negative (−0.11). The monotone increase in within-decile Sharpe ratios at $I = 60$ (from 0.60 at D1 to 0.77 at D10) shows that the CNN *ordering* is still informative, but the gross spread is too small relative to volatility to survive transaction costs.

5.5 Classification metrics

Table 3 shows that AUC rises with horizon, in line with the original paper, which argues that longer look-back windows expose more learnable structure; accuracy at $I = 5$ is *below* the 50% benchmark, which is consistent with a left-skewed label distribution (more days lose money in the weekly rebalance universe) rather than a strictly negative signal: the H–L Sharpe of 1.29 shows that the predicted-probability *ranking* is far more informative than the binary forecast.




Figure 5: Cumulative log returns of the H–L decile spread, net of transaction costs (top row) and gross (bottom row). The $I = 5$ spread is the only one whose net curve still rises persistently; the $I = 20$ and $I = 60$ net curves flatten relative to their gross counterparts because round-trip turnover is in the 84–87% range (Table 2).

Table 2: Equal-weighted decile portfolio statistics (reproduction of Table 1 in the original paper). **Ret** is annualised mean return, **Std** annualised standard deviation, **SR** the resulting Sharpe ratio. “H–L” columns are for the long–short decile spread net of 20 bp round-trip transaction costs.

	D1	D2	D3	D4	D5	D6	D7	D8	D9	D10	H–L
$I = 5$ (weekly rebalance, turnover 0.858)											
Ret	0.07	0.12	0.13	0.13	0.14	0.15	0.16	0.16	0.17	0.21	0.14
Std	0.13	0.16	0.17	0.17	0.18	0.18	0.19	0.19	0.20	0.21	0.11
SR	0.55	0.75	0.77	0.77	0.75	0.79	0.83	0.85	0.89	1.04	1.29
$I = 20$ (monthly rebalance, turnover 0.841)											
Ret	0.10	0.12	0.13	0.12	0.11	0.12	0.11	0.11	0.11	0.11	0.01
Std	0.17	0.18	0.18	0.18	0.18	0.18	0.17	0.17	0.17	0.17	0.08
SR	0.59	0.63	0.72	0.65	0.63	0.65	0.65	0.66	0.68	0.68	0.16
$I = 60$ (quarterly rebalance, turnover 0.867)											
Ret	0.13	0.13	0.13	0.14	0.13	0.12	0.13	0.13	0.12	0.12	−0.01
Std	0.22	0.20	0.20	0.19	0.20	0.19	0.18	0.18	0.17	0.16	0.11
SR	0.60	0.65	0.67	0.71	0.65	0.67	0.69	0.73	0.73	0.77	−0.11

5.6 Fama–French five-factor + momentum spanning

The regression in Table 4 delivers three messages. First, at $I = 5$ the strategy’s annualised alpha is statistically indistinguishable from zero ($+0.28\%$, $t = 0.13$): the high raw Sharpe is largely captured by factor exposures (significant loadings on Mkt–RF, SMB, HML, and an opposite-signed load on CMA). Second, at $I = 20$ and $I = 60$ the alphas are *negative* and significant ($−4.30\%$, $t = −2.23$; $−6.65\%$, $t = −3.77$); the loadings on RMW are positive and large, and at $I = 60$ the strategy also picks up a positive momentum loading. Third, the very high R^2 at $I = 60$ (0.63) suggests that the CNN signal at this horizon is essentially a re-weighting of standard style factors rather than an independent source of premia.

6 Interpretability: input-gradient saliency

Following Simonyan et al. [2014], we compute the absolute pixel-level gradient of the “up” logit with respect to each input image for three high-confidence positive examples drawn from the $I = 60$ test window (Figure 6). Across all three examples the saliency mass is concentrated on the right-most 10–20% of the OHLC panel and on volume bars near sharp price moves. This is consistent with a data-driven momentum reading: the CNN essentially learns to up-weight the most-recent observations, which is also the regime in which classical price-trend signals concentrate their power.

Table 3: Out-of-sample binary classification metrics for the CNN ensemble (2001–2019). N_{test} is the number of rebalance-day observations scored by the five-seed ensemble.

Horizon	Accuracy	AUC–ROC	N_{test}
$I = 5$	49.45%	0.5103	6,378,408
$I = 20$	52.68%	0.5242	1,441,167
$I = 60$	53.70%	0.5288	453,847

Table 4: Fama–French five-factor plus momentum spanning regressions for the net CNN long–short spread. The dependent variable is $r_t^{\text{CNN,LS}} - r_t^f$ regressed on {Mkt–RF, SMB, HML, RMW, CMA, Mom} at the matched rebalance frequency. Alphas are annualised; t -statistics are in parentheses.

	$I = 5$	$I = 20$	$I = 60$
α (annualised, %)	+0.28 (0.13)	−4.30 (−2.23)	−6.65 (−3.77)
$\beta_{\text{Mkt–RF}}$	0.29 (14.24)	0.10 (2.21)	0.08 (1.24)
β_{SMB}	0.14 (3.81)	0.01 (0.11)	−0.14 (−1.24)
β_{HML}	0.13 (3.57)	0.02 (0.27)	0.23 (2.21)
β_{RMW}	0.02 (0.41)	0.32 (3.48)	0.69 (5.24)
β_{CMA}	−0.25 (−4.20)	−0.11 (−0.98)	−0.29 (−2.01)
β_{Mom}	−0.03 (−1.51)	−0.00 (−0.05)	0.23 (3.86)
R^2	0.34	0.07	0.63
N periods	989	226	74

7 Discussion and conclusion

We recover the qualitative shape of the original paper at $I = 5$ – a monotone decile ranking, a high gross-of-cost H–L spread that survives transaction costs into a 1.29 net Sharpe ratio, and a decile-by-decile dominance over momentum and moving-average cross. At $I = 20$ and $I = 60$ the qualitative ranking is preserved (Table 2, Figure 3), but the H–L spread is too thin to survive costs, and the Fama–French six-factor regression delivers a statistically significant *negative* alpha at both horizons.

Several factors plausibly account for numerical differences relative to the original paper. (i) Our CRSP-style extract, liquidity filters, and ETF-based calendar construction need not coincide exactly with the original paper’s. (ii) We aggregate five randomly seeded models; the original paper’s headline results pool a larger ensemble, which can stabilise probability estimates. (iii) Small implementation choices—chart rasterisation, float rounding, and how turnover maps into proportional costs—can move Sharpe ratios and alphas when gross spreads are thin.

The saliency analysis (Figure 6) shows that the $I = 60$ CNN learns a momentum-like reading rather than a more exotic pattern; this is consistent with a model whose alpha is largely spanned by RMW and Mom. Promising directions for future work include regime-conditioned heads, industry- or size-neutralised label construction, and a mixture-of-experts ensemble that explicitly allocates capacity across horizons.

References

- Mark M. Carhart. On persistence in mutual fund performance. *Journal of Finance*, 52(1):57–82, 1997.
- Eugene F. Fama and Kenneth R. French. Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics*, 33(1):3–56, 1993.
- Eugene F. Fama and Kenneth R. French. A five-factor asset pricing model. *Journal of Financial Economics*, 116(1):1–22, 2015.
- Shihao Gu, Bryan T. Kelly, and Dacheng Xiu. Empirical asset pricing via machine learning. *Review of Financial Studies*, 33(5):2223–2273, 2020.




Figure 6: Input-gradient saliency for three high-confidence positive predictions of the $I = 60$ ensemble. Left column: the input chart (inverted to a white background for readability). Right column: the same image with the gradient magnitude $|\partial f_{\text{up}}/\partial x|$ overlaid in red. The network’s attention concentrates on the right-most third of each window (i.e. the most-recent price action) and on volume bars near regime changes.

Narasimhan Jegadeesh and Sheridan Titman. Returns to buying winners and selling losers: Implications for stock market efficiency. *Journal of Finance*, 48(1):65–91, 1993.

Jingwen Jiang, Bryan T. Kelly, and Dacheng Xiu. (re-)imag(in)ing price trends. *Journal of Finance*, 78(6):3193–3249, 2023.

Bryan T. Kelly, Seth Pruitt, and Yinan Su. Characteristics are covariances: A unified model of risk and return. *Journal of Financial Economics*, 134(3):501–524, 2019.

Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.

Andrew W. Lo, Harry Mamaysky, and Jiang Wang. Foundations of technical analysis: Computational algorithms, statistical inference, and empirical implementation. *Journal of Finance*, 55(4): 1705–1765, 2000.

Omer Berat Sezer and Ahmet Murat Ozbayoglu. Algorithmic financial trading with deep convolutional neural networks: Time series to image conversion approach. *Applied Soft Computing*, 70: 525–538, 2018.

Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. Deep inside convolutional networks: Visualising image classification models and saliency maps. In *International Conference on Learning Representations (ICLR) Workshop*, 2014.

A Reproduction artefact map

Every numerical claim in this paper traces back to a file in `final_submit_version/outputs/`:

- Figure 1: produced by `scripts/pipeline_split/export_paper_assets.py`.
- Figure 3, Figure 4: PNGs in `pipeline_runs/figures/`; decile-level CSVs in `pipeline_runs/tables/fig{6,7_simple}_I{5,20,60}_{cnn,ma,mom,reversal}.csv`.
- Figure 5: PNGs in `ppt_images/ppt15_cumulative_returns{,_gross}_L{5,20,60}.png`.
- Table 2: `pipeline_runs/tables/table1_I{5,20,60}.csv`.
- Table 3: `ppt_images/ppt16_classification_metrics.csv` plus shape of `cache/preds_I{5,20,60}.npy`.
- Table 4: produced by `scripts/pipeline_split/export_ff_regression.py`; intermediate file `pipeline_runs/tables/ff_regression_combined.csv`.
- Figure 6: produced by `scripts/pipeline_split/export_saliency_map.py` using the $I = 60$ checkpoint at seed 42.

B Hyperparameters and per-block tensor shapes

Table 5 reports the per-block configuration used in the three CNN backbones; all blocks share the Conv→BatchNorm→LeakyReLU(0.01)→MaxPool(2, 1) pattern.

Table 5: CNN backbone configuration by horizon. “–” denotes a convolutional block not used for that horizon. FC input dimensions match Figure 2.

Block	$I = 5$	$I = 20$	$I = 60$
Block 1 channels	$1 \rightarrow 64$	$1 \rightarrow 64$	$1 \rightarrow 64$
Block 1 stride	(1,1)	(3,1)	(3,1)
Block 1 dilation	(1,1)	(2,1)	(3,1)
Block 2 channels	$64 \rightarrow 128$	$64 \rightarrow 128$	$64 \rightarrow 128$
Block 3 channels	–	$128 \rightarrow 256$	$128 \rightarrow 256$
Block 4 channels	–	–	$256 \rightarrow 512$
FC input dim	15,360	46,080	184,320
Approx. params	~ 150 k	~ 700 k	~ 3.0 M

C Computational footprint

The cleaned daily panel occupies approximately 11.1 GB. Per-seed model checkpoints are approximately 0.6 MB ($I = 5$), 2.8 MB ($I = 20$), and 11.8 MB ($I = 60$); five seeds per horizon give a total of about 76 MB. Training and ensemble inference were performed on a single CUDA GPU (with an Apple-silicon fallback for development); end-to-end execution of the pipeline takes on the order of a few wall-clock hours once the cleaned panel has been constructed.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Section 1 state precisely what we reproduce, with horizon-specific Sharpe ratios, AUCs, and Fama–French alphas; every numerical claim is traceable to a CSV listed in Appendix A.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.

- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 catalogues divergences from the original paper at $I = 20$ and $I = 60$, attributing them to post-publication signal decay, ensemble size, transaction-cost handling, and volume-bar normalisation.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: This is an empirical reproduction study; it contains no novel theoretical results requiring proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.

- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 3 specifies the universe and image construction; Section 4 and Appendix B detail the architecture, optimiser, learning rate, batch size, dropout, and the five training seeds. Appendix A describes the reproducibility protocol used to regenerate the reported tables and figures.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The replication materials include the code, notebooks, trained checkpoints (five seeds per horizon), prediction arrays, and post-aggregation tables needed to reproduce the main results. The underlying daily equity panel is sourced from CRSP under the standard institutional licence; the relevant fields and filters are described in Section 3.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The training recipe (SGD, learning rate 10^{-5} , batch size 128, dropout 0.5, 50-epoch budget with validation-loss early stopping, five seeds) is given in Section 4; per-block tensor shapes are in Table 5; data splits are specified at the end of Section 3.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Table 4 reports t -statistics on every Fama–French factor loading and on the annualised alpha, computed under the OLS-homoskedastic standard errors returned by `statsmodels`. Statistical significance of factor exposures and alphas is therefore explicitly reported.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix C reports the on-disk footprint of the cleaned cache, per-horizon checkpoint sizes, and indicative training time. Training and inference were carried out on a single CUDA GPU with an Apple-silicon MPS fallback for local development.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The project uses publicly described, institutionally licensed CRSP-style equity data only; it does not involve human subjects, scraped personal data, or models with foreseeable safety risks.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The paper is a reproduction of a published empirical asset-pricing study using publicly studied CRSP data; we identify no direct path to material societal benefit or harm beyond what is already discussed in the underlying literature on quantitative trading.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.

- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The released artefacts are small CNN classifiers for next-period U.S. equity return direction; they pose no safety risks comparable to large generative models or scraped image datasets.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The original paper [Jiang et al., 2023] is cited throughout; the daily equity panel is sourced from CRSP under the standard university subscription terms; the Fama–French factors are obtained via the public Kenneth French Data Library through `pandas_datareader`. PyTorch, NumPy, pandas, and matplotlib are used under their respective open-source licences.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released assets are organised by empirical stage, with separate components for data construction, image generation, model training, prediction, portfolio aggregation, factor regressions, and saliency analysis. The trained checkpoints, prediction arrays, and post-aggregation tables are documented alongside the replication materials.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This study does not involve crowdsourcing or human subjects of any kind.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This study does not involve human subjects, so IRB review is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: No large language model is part of the core methodology of this reproduction; the CNN classifier is the only learned component, and any LLM-based assistance was limited to writing and editing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.